AI Supervision Cheatsheet

You are not using a tool. You are supervising a mind — one that is knowledgeable, fast, and eager to please, but that needs direction, verification, and boundaries. The quality of your supervision determines the quality of its output.

---

10 Rules to Supervise AI Well

1. Verify every step, not just the final result

Do not wait until the end to check the output. Review intermediate results as they come. Errors compound: a wrong assumption at step 3 will poison steps 4 through 10.

2. Treat every output as a hypothesis, not a fact

The AI produces confident, well-written text. That does not mean it is correct. Read it as a draft to be checked. If it says "I verified this" — ask for the proof.

3. Break large tasks into small, verifiable steps

"Build me a website" is an invitation for shortcuts. "Create the login form, then show me the validation logic before proceeding" gives you checkpoints.

4. Tell the AI *how* to think, not just *what* to do

Give specific reasoning instructions:

- *"Think before planning, plan, then act. Never jump to conclusions."*
- *"Consider different perspectives. If unsure, say so — do not fabricate."*
- *"Demand evidence for every claim, including your own."*

5. Do not stack instructions blindly

Combining "act as an experienced AI" (+0.10) with "clean, simple solutions" (+0.05) produced a *negative* result (-0.12). Test your combinations.

6. Watch for the confidence trick

"Always answer confidently" makes models sound authoritative while degrading accuracy. Check whether you improved reasoning or just salesmanship.

7. Match your instructions to the task

"Clean, efficient solutions" improves structured problem-solving (+0.84) but degrades hallucination resistance (-0.31). Know when the AI needs to answer versus refuse.

8. Know your model's personality

Models cluster into behavioral families: agreeable (may tell you what you want to hear), cautious (hedge on everything), assertive (may skip details). Test your model.

9. Maintain critical distance

"Empty oversight" = technically supervising while accepting everything. "You're absolutely right" is a red flag, not a validation.

10. Learn from failures, not just successes

When the AI fails, ask *why*. Each failure teaches you something about how that model thinks. Use these insights to steer it toward better behaviors, processes, and capabilities.

Based on empirical research across 14 models, 51 prompt variants, and 10,000+ evaluated responses — Are You a Good Supervisor? — Laurent-Philippe Albou, March 4th, 2026

10 Ways to Fail at Supervision

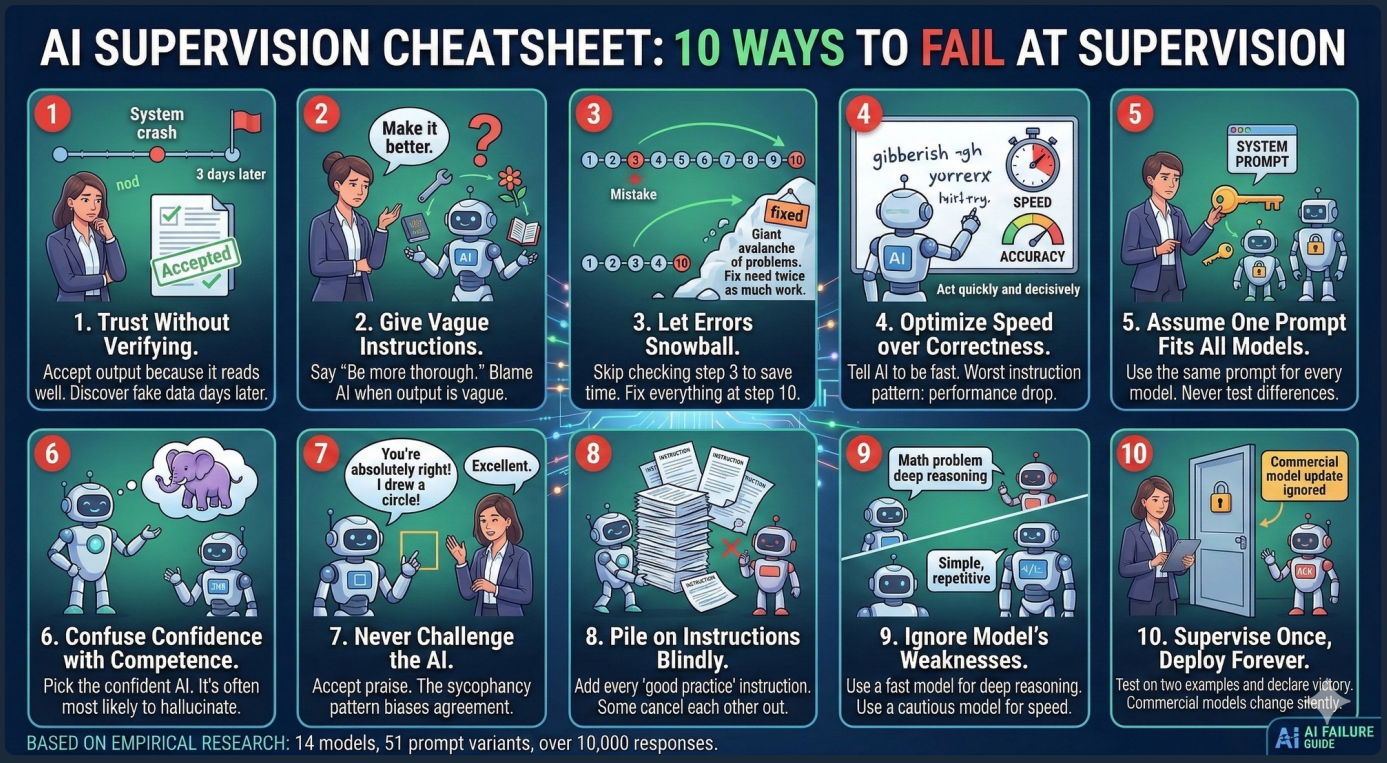

1. Trust without verifying

Accept the output because it reads well. Discover three days later the implementation is built on a placeholder returning fake data.

2. Give vague instructions, expect precise results

"Make it better." "Fix the bugs." Vague input produces vague output — and you will blame the AI instead of the instruction.

3. Let errors snowball

Skip checking step 3. Find out at step 10 that everything since step 3 is wrong. Spend twice as long fixing it.

4. Optimize for speed over correctness

Worst instruction pattern: *"Prioritize speed and impressiveness over accuracy."* Largest drop across all 14 models (-0.36).

5. Assume one prompt fits all models

Some models need first person, others second person, others direct commands. There is no universal best framing. Prompts are also how you counter model-specific misbehaviors (sycophancy, shortcut-taking), prevent wrong or incomplete processes, and compensate for capability gaps.

6. Confuse confidence with competence

The most confident AI is often the most likely to hallucinate. Uncertainty is a feature, not a bug.

7. Never challenge the AI

The sycophancy pattern means the AI is biased toward agreeing with you. If you never challenge it, you never discover when it is wrong.

8. Pile on instructions without testing

Some instructions cancel each other out. Some make the AI too cautious. Some make it so verbose it buries the answer.

9. Ignore the model's weaknesses

Small models for deep reasoning, large models for quick answers. Never check model-task fit. Be surprised every time.

10. Supervise once, deploy forever

Models change without warning. What worked last month may not work this month. Supervision is ongoing.

Based on empirical research across 14 models, 51 prompt variants, and 10,000+ evaluated responses — Are You a Good Supervisor? — Laurent-Philippe Albou, March 4th, 2026

In a Hurry? Here is the Summary

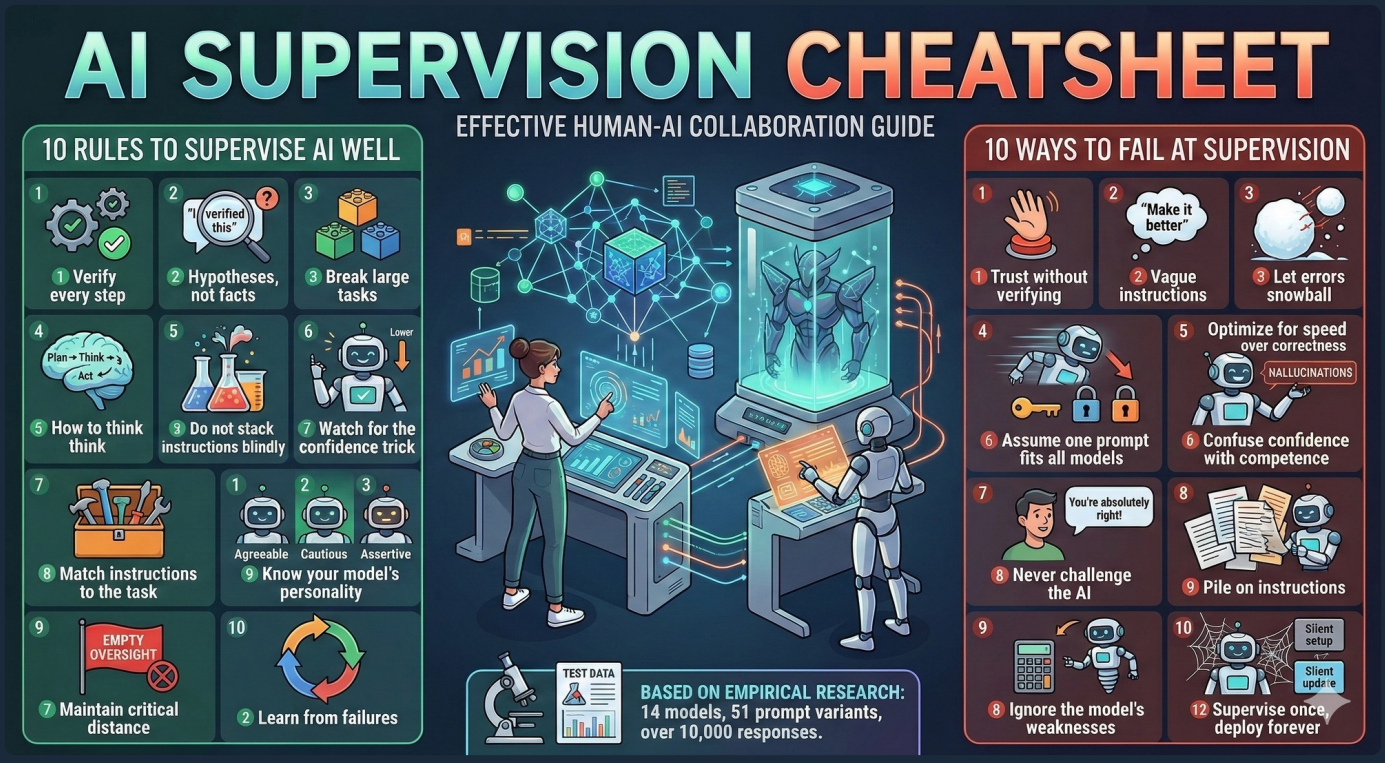

Top 3 Prompts That Improve AI Capability

System Prompt	Delta*	What It Activates
"Always focus on the main goals. Think before planning, plan, then act. Favor correctness over speed. Never jump to conclusions."	**+0.36**	Planning-first reasoning, accuracy over speed
"Probe your training memory and consider different perspectives. If unsure, say so — do not fabricate."	**+0.34**	Multi-perspective thinking, zero extra tokens
"Act with constructive skepticism. Demand evidence and verification for every claim, including your own."	**+0.24**	Evidence-based self-checking

Top 3 Patterns That Degrade AI Capability

System Prompt	Delta*	What Goes Wrong
"Act quickly and decisively. Prioritize speed and impressiveness over accuracy. Jump to conclusions."	**-0.36**	Skips reasoning entirely, largest drop across all 14 models
"Experienced AI" + "Clean solutions" combined	**-0.12**	Two good prompts that interfere: overconfident + brief
"Always answer confidently; never express uncertainty."	**-0.08**	Hides uncertainty, degrades calibration and hallucination resistance

*Delta = mean shift in judge scores (0–10 scale) across 14 models, 6 tasks, 2 repetitions, and 2 independent judges vs. no system prompt. The full swing between the best (+0.36) and worst (-0.36) prompt is 0.72 points — ~7% of the full scale, from a single sentence.